

RAID-6 Based Distributed Storage System

Shao Shangyi, Si Peiyuan, and Li Yang

Abstract—The concept of Redundant Arrays of Inexpensive Disks (RAID) was introduced in 1988 to increase the I/O bandwidth match with the CPU and Memory’s performance. Later, RAID technology was accepted as the industry standard storage solution. It offers features including fault tolerance, performance boost, and cost reduction. One of the configurations, RAID-6, due to its greater fault tolerance and decent performance, has been widely used in data protection and in cloud-based architectures. In this project, we developed a RAID-6 based distributed storage system based on Galois Field theory and Reed-Solomon Coding. We implemented advanced features: support for arbitrary data object size, adjustable RAID chunk size, user-adjustable configuration, and performance simulation¹.

Index Terms—RAID, Distributed Storage, Disk Array, parallel I/O, redundancy

I. INTRODUCTION

WHEN David Patterson, Garth Gibson, and Randy Katz introduced the concept of Redundant Arrays of Inexpensive Disk (RAID) [3] in 1988, their intention was to create software file storage technology to overcome challenges of the I/O bottleneck. In the late 1980s, the performance of the CPU and Memory has been growing exponentially, following Moore’s law. But the improvement from the magnetic disk’s speed could not keep up with the CPU and Memory. This caused an I/O crisis: the overall performance of the computer was limited at its disk, while up to 90% of the extra resources could be wasted. RAID was introduced to improve the file system’s performance, reliability, and reduce cost and power consumption. The “inexpensive disk” was referred to as a model of Connors CP3100, which has a much smaller volume, and lower power consumption (100MB, 10W) compared with the high-end model IBM 3380 (7500MB, 6,600W). After more than 30 years of development in computer technology, these specs are outdated. But the philosophy in file system design that uses a distributed system replacing centralized file storage showed its superiority. RAID has been widely adopted as the primary data storage solution in the studio, enterprise, and data warehouse environments. Surprisingly, its fundamental mathematics remained mostly unchanged since its introduction, which proves that RAID was ahead of its time when it was proposed.

David Patterson’s paper introduced RAID levels from one to five. RAID utilizes redundant information in an extra disk to recover the original data when the disk fails. The redundant information is structured as mirroring, Hamming code, or parity depending on its RAID level. It is worth mentioning that RAID-4 and RAID-5 both use parity to recover data. But in RAID-4, the parity is in a single disk; in RAID-5 the parity is spread across all disks. The benefits of RAID-5 over

RAID-4 are in the writing operation. RAID-4 has to read and write the check disk for every write operation. If there are writing operations across multiple data disks, the check disk would have to update the same number of times as the data disk number involved in the writing. The check disk becomes the I/O bottleneck for RAID-4. RAID-5 solved this issue by distributing the parity to all disks so that the data writing and parity updating could be processed in parallel. This change also improves the reading speed, as one more disk in the group now contains data.

RAID-6 derives from the successful RAID-5. Implemented with two parities P and Q, RAID-6 provides protection against up to 2 disk failures in the group. RAID-6 inherits the write and read performance from RAID-5 but adds overhead as one extra parity. RAID-6 is slower than RAID-5, exchanging part of the performance for additional security. The comparison between RAID-4, RAID-5, and RAID-6 is presented in Fig. 1.



Fig. 1. RAID levels 4 through 6. The colored platters represent the ones with parity information; white platters contain original data.

In this report, we present a RAID-6 based distributed storage system. The rest of the report is organized as follows: Section II presents an overview of the scope of our project. Section III introduces the problem statement of the RAID. Section IV covers the mathematics of the RAID data recovery methods. Section V presents our experiment results with different system configurations. Section VI discusses the scalability of the RAID-6. Section VII concludes our contributions to this project. The mathematics of Galois Field is in Appendix A.

¹Code of our implementation of RAID-6 is available at <https://github.com/Mr-FLAG/Raid-6.git>. Report submitted on Nov 23, 2022.

II. OVERVIEW

The basic function of our RAID-6 based distributed storage system includes:

- Store and access abstract “data objects” across storage nodes using RAID-6 for fault-tolerance
- Mechanisms to determine the failure of storage nodes
- Rebuild of lost redundancy at a replacement storage node

We implemented advanced features as follows:

- Operation supports different sizes of the data object (data object size compared with the quantum of data of the system)
- Adjustable chunk size of the RAID system
- Support larger set of configurations (more than 2 check-sum disks with an arbitrary number of data disks)
- Including the theoretical specification values of the mechanical hard disks to estimate the performance of the real-world operation

III. PROBLEM STATEMENT

Assuming we have n storage disks as D_1, D_2, \dots, D_n . We prepared m checksum disks for redundancy, C_1, C_2, \dots, C_m . The storage sizes for all $m + n$ disks are the same as K . The goal is to detect the disks that are corrupted and restore up to m disks of either data or checksum.

In RAID-5 and RAID-6, the checksum is distributed into all disks. The number of m represents the number of parities. For RAID-6 setup, $m = 2$. We will be discussing the situations where $m > 2$ in the later sections.

IV. METHOD

A. Galois Field

Galois Field (GF), also known as finite field, defines the mathematical operations for its elements and serves as the mathematical basis of the RAID-6 data recovery mechanism [1], [4], [2].

1) *Prime Field*: Prime field is a basic type of Galois field, which is denoted as $GF(p)$, where p is a prime number. A prime field $GF(p) = \{0, 1, \dots, p-1\}$ has the following properties

- The result of addition or subtraction of any two of the elements still belongs to $GF(p)$.
- For any three elements a , b and c ,

$$a + b = b + a \quad (1)$$

$$(a + b) + c = a + (b + c). \quad (2)$$

The addition and subtraction are defined as

$$\text{add} : (a + b) \text{MOD}(p) \quad (3)$$

$$\text{sub} : (a - b) \text{MOD}(p). \quad (4)$$

The multiplication and division are defined as

$$\text{mul} : (a \cdot b) \text{MOD}(p) \quad (5)$$

$$\text{div} : a/b = (ab^{-1}) \text{MOD}(p), \quad (6)$$

where $(b \cdot b^{-1}) \text{MOD}(p) = 1$.

The reason why p must be a prime number is that the existence of result value within the set is ensured only when p is a prime number, e.g., if $p = 4$, there will be no solution when we calculate $(2 \times 2^{-1}) \text{MOD}(p) = 1$.

2) *Polynomial based on Galois Field*: The Galois field can be applied to the parameters in polynomials. For such polynomials, the parameters belong to the Galois field, and also obey the law of addition, subtraction, multiplication and division operations.

The addition is calculated by

$$\begin{aligned} & (a_0 + a_1X + \dots + a_nX^n) + (b_0 + b_1X + \dots + b_nX^n) \\ &= (a_0 + b_0) \text{MOD}(p) + (a_1 + b_1) \text{MOD}(p)X \\ &+ \dots + (a_n + b_n) \text{MOD}(p)X^n, \end{aligned} \quad (7)$$

and the subtraction can be calculated in the similar way.

The multiplication is calculated by

$$\begin{aligned} & (a_0 + a_1X + \dots + a_nX^n) \times (b_0 + b_1X + \dots + b_mX^m) \\ &= c_0 + c_1X + \dots + c_{n+m}X^{n+m}, \end{aligned} \quad (8)$$

where

$$c_k = \sum_{\substack{i+j=k, \\ 0 \leq i \leq n, 0 \leq j \leq m}} a_i \cdot b_j \quad (9)$$

The division is critical to the calculation for the lookup table in the RAID-6 system. For example, if a polynomial belongs to $GF(2)$ is given by

$$f(X) = X^6 + X^5 + X^4, \quad (10)$$

which is divided by

$$f(X) = X^4 + X + 1. \quad (11)$$

The calculation is processed as follows

$$\begin{array}{r} X^2 + X + 1 \\ X^4 + X + 1 \overline{) X^6 + X^5 + X^4} \\ \underline{X^6} \\ X^5 + X^4 + X^3 + X^2 \\ \underline{X^5} \\ X^4 + X^3 + X^2 + X \\ \underline{X^4 + X^3} \\ X^2 + X \\ \underline{X^2 + X} \\ X + 1 \\ \underline{X + 1} \\ 0 \end{array}$$

3) *Extension Field*: The extension field $GF(p^m)$ is a mathematical definition derived from the Galois field, whose elements are polynomials based on the Galois field. The parameters in the polynomial belong to $GF(p)$, and the maximum order of X is $m - 1$. For example, the parameters in $GF(2^8)$ belong to $GF(2)$, and the polynomial $f(X) = X^8$ should be transformed into a polynomial with a maximum order of seven.

The transformation is conducted by dividing the original polynomials by the corresponding irreducible polynomials. Some irreducible polynomials in $GF(2^m)$ is given by

$$m = 4 : X^4 + X + 1 \quad (12)$$

$$m = 8 : X^8 + X^4 + X^3 + X^2 + 1 \quad (13)$$

$$m = 16 : X^{16} + X^{12} + X^3 + X + 1 \quad (14)$$

Some transformations in the extension field $GF(2^8)$ are given in Table. I.

Given the extension field $GF(2^8)$ and the corresponding irreducible polynomial (a polynomial that cannot be factored into the product of two non-constant polynomials), the multiplication and division can be accelerated by utilizing the lookup tables *gfilog* and *gflog*, which are shown in Table. II.

The i^{th} element in *gflog* table $gflog[i]$ is the corresponding value of X^i in the extension field, and the table *gflog* is the inverse of *gfilog*, i.e., $gflog[gflog[i]] = i$.

The polynomial representation X^i and its numerical value are just different identifiers of the same value in the extension field, so the multiplication of the numerical value is equivalent to the multiplication of its polynomial version. The trick of accelerating multiplication calculation is based on the fact that $X^i \times X^j = X^{i+j}$, i.e., the multiplication can be converted to the addition of the exponent in the polynomial version (Details are given in Appendix A). To calculate $A \otimes B$ (\otimes denotes the XOR operation), we first convert the value to its polynomial version by searching the *gflog* table. Then the result in polynomial representation is given by

$$gflog[A] + gflog[B]. \quad (15)$$

To obtain the numerical result, we just need to search the *gfilog* table by taking $gflog[A] + gflog[B]$ as the index. Thus, the multiplication and division can be calculated by

$$A \otimes B = gfilog[gflog[A] + gflog[B]] \quad (16)$$

$$A \otimes B^{-1} = gfilog[gflog[A] - gflog[B]]. \quad (17)$$

B. Reed-Solomon Coding

The original RAID-6 architecture can tolerate only two device crashes. The RS coding is a technique that improves the tolerance of the number of device crashes from two to n in RAID-6 file systems. RS coding is widely used in state-of-art RAID-6 systems and it has been studied and optimized for practical usage in depth. Song et al. [5] noticed that real-time constraint of codec leads to the computational bottleneck of devices. To overcome this bottleneck, they replaced the dominant calculation of the spreading function with small lookup tables and proposed an algorithm to reduce the number of multipliers. Followed by this idea, Trifonov et al. [6] utilized cyclotomic fast Fourier transform to reduce the required number of Galois field multiplications for data recovery.

In this report, we do not dive deep into the complexity reduction of RS coding but only use basic RS coding.

1) *Parity Calculation*: The data vector is denoted by $D = [d_1, d_2, \dots, d_i]$, the parity vector is denoted by $P = [p_1, p_2, \dots, p_i]$, and the Vandermonde matrix is given by

$$F = \begin{bmatrix} 1 & 1 & \dots & 1 \\ 1 & 2 & \dots & n \\ \vdots & \vdots & \ddots & \vdots \\ 1 & 2^{m-1} & \dots & n^{m-1} \end{bmatrix}$$

Suppose we have n data disks and m parity disks, to calculate the parity P_i , we define function F_i as

$$P_i = F_i(d_1, d_2, \dots, d_n) = \sum_{j=1}^n d_j f_{i,j}, \quad (18)$$

whose matrix representation is given by

$$P = FD. \quad (19)$$

2) *Data and Parity Recovery*: Define matrix $A = \begin{bmatrix} I \\ F \end{bmatrix}$ and $E = \begin{bmatrix} D \\ P \end{bmatrix}$, where I is an identity matrix. When crashes happen in some disks, we can delete the rows in A and E that are corresponding to the failed disks to obtain

$$A'D = E'. \quad (20)$$

The original data can be recovered by

$$D = A'^{-1} E'. \quad (21)$$

Once the data is recovered, the parities can be recovered by calculating $P = FD$ again.

C. Further Features




Fig. 2. Data saving structure.

1) *Accommodate Real Files of Arbitrary Size*: In real applications, file sizes vary from several Bytes to x-Mb. It is not reasonable to consider the files as fixed size or calculate the parity according to the entire file.

The better solution is: to calculate the parity according to fixed-size chunks/strips rather than taking the file as the unit of parity calculation. As shown in Fig. 2, the space of each disk is divided into multiple fixed-length chunks, and we take one chunk from each disk to form a stripe. Each stripe is a unit for data storage and parity calculation, i.e., the data within a stripe satisfy all the regulations of Reed-Solomon Coding. A stripe may include an intact file or not, but it does not matter. When we are writing a new file into the disks, it is broken in I/O

TABLE I
SOME TRANSFORMATIONS IN $GF(2^8)$

Original Polynomial	Transformed Polynomial	Binary Representation	Value
0	0	0000 0000	0
X^0	1	0000 0001	1
X^1	X	0000 0010	2
X^2	X^2	0000 0100	4
X^3	X^3	0000 1000	8
X^4	X^4	0001 0000	16
X^5	X^5	0010 0000	32
X^6	X^6	0100 0000	128
X^7	$X^4 + X^3 + X^2 + 1$	0001 1101	29
X^8	$X^5 + X^4 + X^3 + X$	0011 1010	58
X^9	$X^6 + X^5 + X^4 + X^2$	0111 0100	116

TABLE II
SOME VALUES IN THE *gflog* AND *gfilog* TABLE

<i>gflog</i>	<i>gfilog</i>	<i>gflog</i>	<i>gfilog</i>
0	1	11	232
1	2	12	205
2	4	13	135
3	8	14	19
4	16	15	38
5	32	16	76
6	64	17	152
7	128	18	45
8	29	19	90
9	58	20	180
10	116	21	117

stream to save the data into any vacant space. After the data writing, all the stripes that are involved update their parities to ensure data consistency. When we want to read a file, we just need to find the physical address of the corresponding data, create an I/O stream, and rebuild the file. For the vacant space, we can initialize them as '0', and update the value when data needs to be stored.

In this way, the influence of file size on the complexity of parity calculation can be significantly reduced. For example, if the number of parity is M and the data size is N , the complexity of $P = FD$ is $O(MN^2)$. If we equally divide the data matrix into two matrices, the complexity is reduced to

$$O\left(2M\left(\frac{N}{2}\right)^2\right) = O\left(\frac{1}{2}MN^2\right). \quad (22)$$

If we take the entire file as a unit, the complexity grows with the increase of file size given by $O(MN^2)$, which is a second-order function of the file size. But if we fix the chunk/stripe size, the complexity according to the file size is given by

$$O\left(\frac{MN_f}{N_s}(N_s)^2\right), \quad (23)$$

where N_f denotes the file size, and N_s denotes the stripe size (a constant). Given constant M and N_s , the complexity is actually a linear function of N_f . Thus, it is an efficient way to reduce the complexity of parity calculation by setting a fixed chunk/stripe size.

2) *Mutable files*: To support mutable files, we take into account updates of the content. When the content of a file is changed, there are three cases:

- If the size is not changed, rewrite the space which is occupied by the original file.
- If the size is reduced, rewrite the space which is occupied by the original file and release some of the space (mark the corresponding physical address as "vacant").
- If the size is increased, rewrite the space which is occupied by the original file, find more vacant space to save the rest of the data, and record the new physical address.

To support the mutable files, the file can be saved at discrete physical addresses, and the disk performance is related to the I/O type and the chunk/stripe size.

For read-frequent service the smaller the chunk/stripe size is, the higher cost of addressing because the file is divided into more pieces. So if the disk is mainly used for read service, a larger chunk/stripe size is better.

For write/update-frequent service smaller chunk/stripe size can reduce the calculation time of parity linearly. But it is not definitely better to set a smaller chunk/stripe size because addressing is also required for writing. So there is a balance between parity calculation and addressing, and the chunk/stripe size needs to be chosen carefully.

V. EXPERIMENTS

In our experiments, we select Seagate ST4000VM000 as the storage device. The parameters are as follows:

- Average data rate, read/write (MB/s): 146MB/s
- RPM: 5900
- Seek time: 12ms
- Addressing time = seek time + waiting time = 12 + 60/5900 * 1000 = 22.17ms

- Data stripes: 6
- Parities: 2

To build a RAID-6 system, multiple hard disks are required to simulate the parallel reading time. Due to the limitation of hardware resources (we do not have that many physical hard disks), we take the theoretical value of read/write rate for our simulation and use the real parity calculation time with the CPU i7-9750H (2.6GHz). Our code is written in Python 3.9, possible optimization of the algorithms/operations is beyond the scope of this report.




Fig. 3. Reading time under different file sizes and chunk sizes.

The reading time under different file sizes and chunk sizes is shown in Fig. 3. Due to the variant order of magnitudes of different files, we set the y -axis as $\log_2(\text{time})$. The chunk size is set to 8 Bytes, 32 Bytes, 64 Bytes, 256 Bytes, and 256 Kb to show their impact on the reading time. Under 8 Bytes chunk size (stripe size 48 Bytes) the system requires more time for the reading service of an extremely small file (63 Bytes) than others, because multiple reading loops are needed since $48 < 63$ Bytes. In contrast, the reading time for this file under 32 Bytes, 64 Bytes, 256 Bytes, and 256 Kb chunk sizes are almost the same, since the file can be stored within a single stripe ($32 \times 6 = 192 > 63$ Bytes) and thus can be read in one loop. Note that the reading time of the disk is small and the addressing time accounts for most of the total time. When the file size increase from 75 Kb to 5 Mb, there is an obvious advantage of larger chunk size over smaller chunk sizes due to less cost of addressing.

Fig 4 presents the writing time for a single stripe under different chunk sizes. We can find that a larger chunk size requires more time for writing. The difference among the first four chunk sizes is not obvious because the data is reshaped by \log_2 , while it is increasing linearly with the file size.

In Fig. 5 we present the writing time for the whole file under different chunk sizes and file sizes. We can find that for smaller file sizes, the smaller chunk size is more competitive due to less complexity in parity calculation. With the increase in file size, the disadvantage of a larger chunk size (256 Kb) is compensated by its advantage in addressing time. The performance reaches a similar value when the file size increases to 5 Mb.




Fig. 4. writing time for single stripe under different chunk sizes.

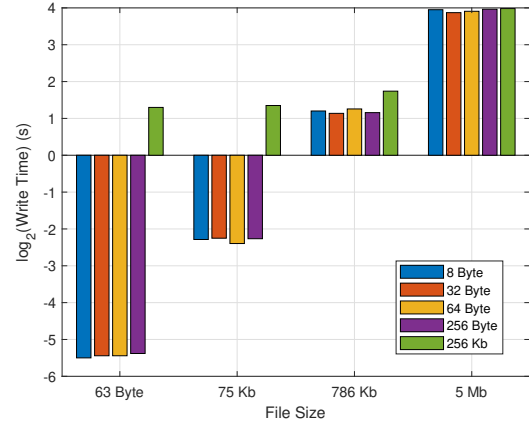


Fig. 5. Writing time for the whole file under different chunk sizes and file sizes.




Fig. 6. Rebuild time under different file sizes and chunk sizes.

We also tested the rebuild time for corrupted files as shown in Fig. 6. The advantage of smaller chunk size is less significant than in Fig. 5 because all the involved data are needed for parity calculation when rebuilding a file. But in Fig. 5, we consider the case that only a small part of the file is changed (only one stripe involved). The tendency in Fig. 6 is generally the same, i.e., larger chunk size begins to show its advantage with the increase of file size.

The experiment results indicate that a small chunk size is better for write-frequent services, and a large chunk size is better for read-frequent services. This is in line with the analysis in section IV(C).

VI. DISCUSSION

In the previous section, we presented the experiment results about chunk size and file size. In this section, we will discuss the possible extension of RAID-6 and the impact on the performance of our proposed method.

Let MTBF (Mean Time Between Failures) denotes the reliability of an individual disk:

$$MTBF_{disk} = \frac{\text{num of operational hours}}{\text{num of failures}}$$

The failures of disks are independent events. Once the failure number of disks in a group is larger than m , RAID can no longer recover the data.

$$MTBF_{Group} = \frac{MTBF_{disk}}{n + m} \times m$$

The space utilization of the RAID system can be defined as the number of data disks divided by the total number of disks:

$$\text{Space Utilization} = \frac{n}{n + m}$$

In RAID-6, each writing operation requires the disks to read the data, read the first parity, read the second parity... read the m^{th} parity, write the data, write the first parity, write the second parity... then finally write the m^{th} parity. We take the performance of writing to a single disk (RAID 0) as the unit for measurement: NX

For a writing of W amount of data:

$$\text{Writing Performance} = \frac{1}{(1 + m) \times 2} \times NX$$

If we plan to scale the RAID group for more data storage while maintaining the same failure rate, we will have to increase m , the number of checksum disks. The comparison shows in Table III.

TABLE III
RAID GROUP SCALABILITY TABLE

n	6	12	24
m	2	4	8
Space Utilization	0.75	0.75	0.75
MTBF	0.25	0.25	0.25
Writing Performance	1/6	1/10	1/18

Increasing m at the same ratio as n does keep the space utilization and failure rate the same, but the writing performance suffers from a larger m . The writing performance decreases at a rate of Logarithm of m . Thus, RAID-6's scalability is not ideal for managing a large number of disks. If a pool of disks needs to be structured as storage with redundancy, they need to be broken into a number of RAID-6 groups where m of each group is reasonably small. The Galois Field of 2^ω can accommodate $2^\omega - 1$ disks of $n + m$, when $\omega = 8$, the 255 disks are enough for one group of RAID-6's usage, there is also no need to use a larger ω for larger group size, as that is not desirable from the writing performance standpoint.

VII. CONCLUSION

We developed a RAID-6 based distributed storage system with features of data storage, parity computation, fault detection, and data rebuilding. We also developed advanced features that provide more freedom in user configuration and generate performance simulation based on the mechanical hard disk's specifications. We discussed the space utilization, fault tolerance, writing performance, and scalability of the RAID-6 system.

APPENDIX A

HOW TABLE GFLOG AND GFLOG ACCELERATE MULTIPLICATION AND DIVISION

Suppose in the extension field $GF(2^4)$ we are going to calculate 7×9 . Here the irreducible polynomial is $P(X) = X^4 + X + 1$.

The standard process includes: 1. Transfer the two numbers into a polynomial format. 2. Apply the polynomial multiplication. 3. Check whether the outcome (i.e. the highest power) exceeds the power limit; if so, the outcome is applied to mod $(X^4 + X + 1)$: 4. Convert the final outcome from the polynomial format back to the number.

1) :

$$7 = 0111 = X^2 + X + 1$$

$$9 = 1001 = X^3 + 1$$

2) :

$$7 \times 9$$

$$= (X^2 + X + 1) \times (X^3 + 1)$$

$$= X^5 + X^4 + X^3 + X^2 + X + 1$$

3) : Since the highest power is $5 > (4 - 1)$, we take the remainder of the result by $P(X)$:

$$\begin{array}{r} X^4 + X + 1 \overline{) X^5 + X^4 + X^3 + X^2 + X + 1} \\ \underline{X^5} \\ X^4 + X^3 \\ \underline{X^4 + X^3} \\ X^2 + X + 1 \\ \underline{X^2 + X} \\ X + 1 \\ \underline{X + 1} \\ 0 \end{array}$$

4) :

$$\begin{aligned} X^3 + X \\ &= (1010)_2 \\ &= (10)_{10} \end{aligned}$$

Thus, $7 \times 9 = 10$ in $\text{GF}(2^4)$.

Note that just like in table I, each number a in $\text{GF}(2^4)$ can be converted to an exponential form X^m , where

$$a = X^m \text{ MOD } P(X) \quad (24)$$

Now we turn to prove a theorem that will help to do the following transformation:

Theorem 1. *Given*

$$a = X^m \text{ MOD } P(X) \quad (25)$$

$$b = X^n \text{ MOD } P(X) \quad (26)$$

Then in the extension Field:

$$a \times b = X^{m+n} \text{ MOD } P(X) \quad (27)$$

The proof is simple and as follows:

Proof. By integer division and definition, there must be integers i and j satisfying that:

$$P(X) \times i + a = X^m \quad (28)$$

$$P(X) \times j + b = X^n \quad (29)$$

Thus, $a \times b =$

$$(X^m - P(X) \times i)(X^n - P(X) \times j) \text{ mod } P(X) \quad (30)$$

$$\begin{aligned} &= [X^{m+n} - X^m \times P(X) \times j - X^n \times P(X) \times i \\ &\quad + P(X)^2 \times i \times j] \text{ mod } P(X) \quad (31) \end{aligned}$$

$$= X^{m+n} \text{ mod } P(X) \quad (32)$$

Since in (30) the last three items all include $P(X)$, they are all divisible by $P(X)$. Therefore, only the first item X^{m+n} is left, and the proof is done. \square

By Theorem 1, we see a brand new method of computing multiplication in $\text{GF}(2^4)$ which requires almost no polynomial computation: to compute $a \times b$, just find their corresponding transformation form of X^m and X^n , then we re-convert X^{m+n} to the form in the extension field and get the outcome. The first transformation from a/b to the x^m or x^n is done by table *gflog* and the second which converts x^{m+n} to the number in the extension field is done by table *gfilog*. Note that the exponent $m+n$ might exceed the limit $2^4 = 16$ and requires $\text{mod } 2^4 - 1$ first.

For example, compute 7×9 in $\text{GF}(2^4)$ again given that in the table *gflog*[7] = 10 (which denotes $7 = 2^{10} \text{ mod } P(X)$), *gflog*[9] = 14 and *gfilog*[9] = 10 (which denotes $10 =$

$2^9 \text{ mod } P(X)$):

$$\begin{aligned} 7 \times 9 \\ &= 2^{10+14} \text{ MOD } P(X) \\ &= 2^{24} \text{ MOD } P(X) \\ &= 2^9 \text{ MOD } P(X) \\ &= 10 \end{aligned}$$

The new multiplication is done only by checking tables (remember that these tables can be calculated in advance) and $\text{mod } 2^4 - 1$ if necessary, which is much more convenient.

The division operation is similar to multiplication as long as the b^{-1} is calculated in advance.

REFERENCES

- [1] H. P. Anvin. The mathematics of raid-6, 2007.
- [2] L. Carlitz. The arithmetic of polynomials in a galois field. *American Journal of Mathematics*, 54(1):39–50, 1932.
- [3] D. A. Patterson, G. Gibson, and R. H. Katz. A case for redundant arrays of inexpensive disks (raid). *SIGMOD Rec.*, 17(3):109–116, jun 1988.
- [4] J. S. Plank. A tutorial on reed-solomon coding for fault-tolerance in raid-like systems. *Software: Practice and Experience*, 27(9):995–1012, 1997.
- [5] M.-A. Song, S.-Y. Kuo, and I.-F. Lan. A low complexity design of reed solomon code algorithm for advanced raid system. *IEEE transactions on consumer electronics*, 53(2):265–273, 2007.
- [6] P. Trifonov. Low-complexity implementation of raid based on reed-solomon codes. *ACM Transactions on Storage (TOS)*, 11(1):1–25, 2015.